

Healthcare

October 2, 2024

Project 4 - Health Care

Problem statement: Cardiovascular diseases are the leading cause of death globally. It is therefore necessary to identify the causes and develop a system to predict heart attacks in an effective manner. The data below has the information about the factors that might have an impact on cardiovascular health.

Dataset description:

Variable	Description
Age	Age in years
Sex	1 = male; 0 = female
cp	Chest pain type
trestbps	Resting blood pressure (in mm Hg on admission to the hospital)
chol	Serum cholesterol in mg/dl
fbs	Fasting blood sugar > 120 mg/dl (1 = true; 0 = false)
restecg	Resting electrocardiographic results
thalach	Maximum heart rate achieved
exang	Exercise induced angina (1 = yes; 0 = no)
oldpeak	ST depression induced by exercise relative to rest
slope	Slope of the peak exercise ST segment
ca	Number of major vessels (0-3) colored by fluoroscopy
thal	3 = normal; 6 = fixed defect; 7 = reversible defect
Target	1 or 0

Task to be performed:

1. Preliminary analysis: a. Perform preliminary data inspection and report the findings on the structure of the data, missing values, duplicates, etc. b. Based on these findings, remove duplicates (if any) and treat missing values using an appropriate strategy.
2. Prepare a report about the data explaining the distribution of the disease and the related factors using the steps listed below: a. Get a preliminary statistical summary of the data and explore the measures of central tendencies and spread of the data. b. Identify the data variables which are categorical and describe and explore these variables using the appropriate tools, such as count plot. c. Study the occurrence of CVD across the Age category. d. Study the composition of all patients with respect to the Sex category. e. Study if one can detect heart attacks based on anomalies in the resting blood pressure (trestbps) of a patient. f. Describe the relationship between cholesterol levels and a target variable. g. State what

- relationship exists between peak exercising and the occurrence of a heart attack. h. Check if thalassemia is a major cause of CVD. i. List how the other factors determine the occurrence of CVD. j. Use a pair plot to understand the relationship between all the given variables.
3. Build a baseline model to predict the risk of a heart attack using a logistic regression and random forest and explore the results while using correlation analysis and logistic regression (leveraging standard error and p-values from statsmodels) for feature selection.

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: data = pd.read_excel("data.xlsx")
```

```
[3]: data.head()
```

```
[3]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
[4]: data.shape
```

```
[4]: (303, 14)
```

```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
```

```
5   fbs      303 non-null   int64
6   restecg  303 non-null   int64
7   thalach  303 non-null   int64
8   exang    303 non-null   int64
9   oldpeak  303 non-null   float64
10  slope    303 non-null   int64
11  ca       303 non-null   int64
12  thal     303 non-null   int64
13  target   303 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
[6]: data.dtypes
```

```
[6]: age          int64
sex          int64
cp          int64
trestbps    int64
chol        int64
fbs         int64
restecg     int64
thalach     int64
exang       int64
oldpeak     float64
slope       int64
ca          int64
thal        int64
target      int64
dtype: object
```

```
[7]: # Checking for missing values
data.isnull().sum(axis = 0)
```

```
[7]: age          0
sex          0
cp          0
trestbps    0
chol        0
fbs         0
restecg     0
thalach     0
exang       0
oldpeak     0
slope       0
ca          0
thal        0
target      0
```

dtype: int64

```
[8]: # To check for duplicates
data.duplicated().sum()
```

[8]: 1

```
[9]: data.drop_duplicates()
```

```
[9]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63	1	3	145	233	1	0	150	0	2.3	
1	37	1	2	130	250	0	1	187	0	3.5	
2	41	0	1	130	204	0	0	172	0	1.4	
3	56	1	1	120	236	0	1	178	0	0.8	
4	57	0	0	120	354	0	1	163	1	0.6	
..	
298	57	0	0	140	241	0	1	123	1	0.2	
299	45	1	3	110	264	0	1	132	0	1.2	
300	68	1	0	144	193	1	1	141	0	3.4	
301	57	1	0	130	131	0	1	115	1	1.2	
302	57	0	1	130	236	0	0	174	0	0.0	

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1
..
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

[302 rows x 14 columns]

```
[10]: data.describe()
```

```
[10]:
```

	age	sex	cp	trestbps	chol	fbs	\
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	

max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000
-----	-----------	----------	----------	------------	------------	----------


	restecg	thalach	exang	oldpeak	slope	ca \
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
mean	0.528053	149.646865	0.326733	1.039604	1.399340	0.729373
std	0.525860	22.905161	0.469794	1.161075	0.616226	1.022606
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000
50%	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000
75%	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000
max	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000

	thal	target
count	303.000000	303.000000
mean	2.313531	0.544554
std	0.612277	0.498835
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

We can see that the scale of each feature column is different and varied.

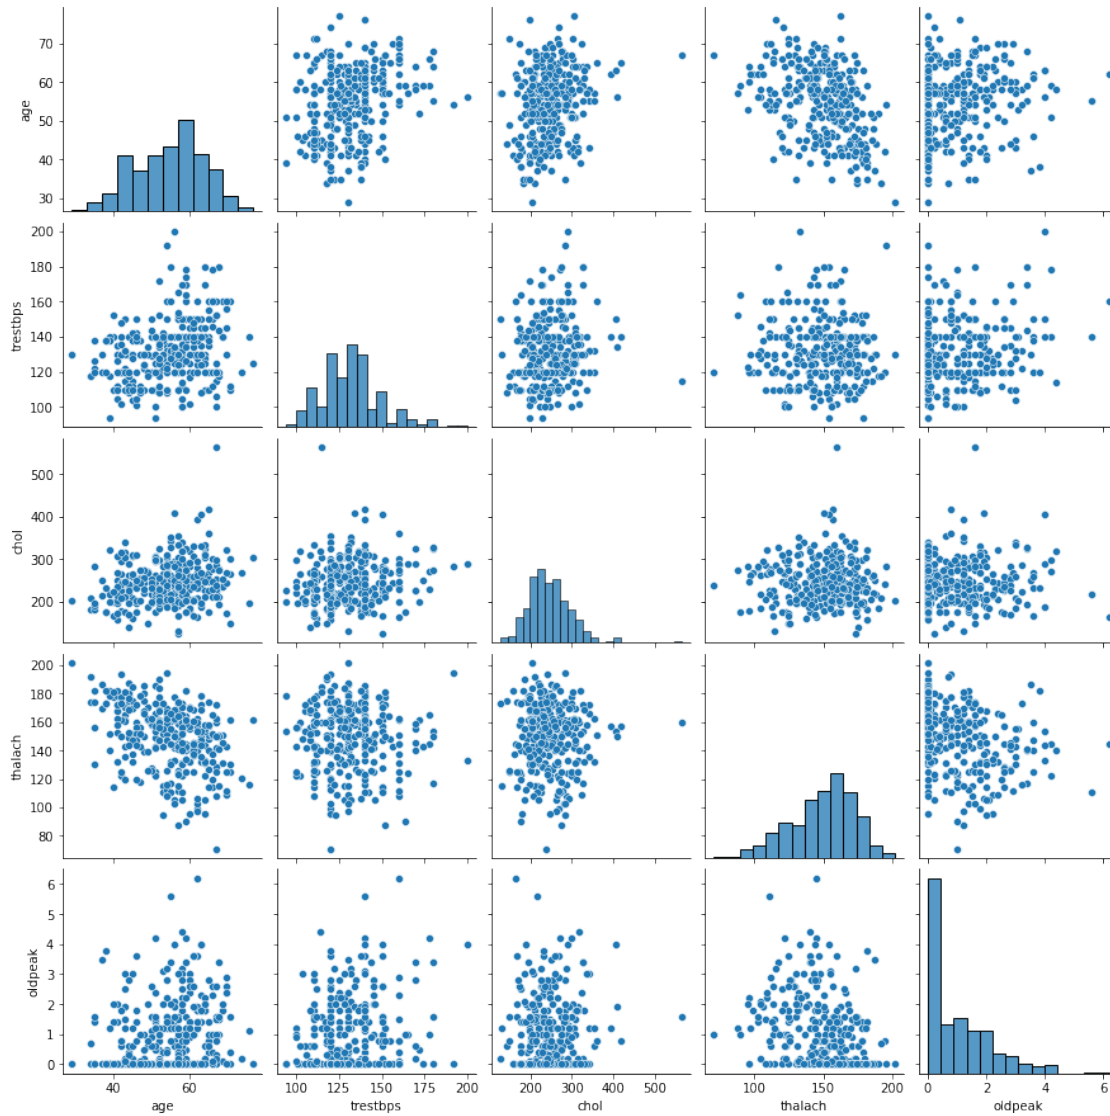
Understanding the Data

```
[11]: # To check the correlation between variables and identify them calculate
      ↪ correlation matrix
corr = data.corr()
plt.subplots(figsize=(15,10))
sns.heatmap(corr, xticklabels=corr.columns, yticklabels=corr.columns,
      ↪ annot=True)
sns.heatmap(corr, xticklabels=corr.columns, yticklabels=corr.columns, annot=True)
plt.show()
```



From the above HeatMap, we can see that *cp* and *thalach* are the features with highest positive correlation whereas *exang*, *oldpeak* and *ca* are negatively correlated. While other features do not hold much correlation with the response variable “target”.


```
[12]: # Using pairplots to see the continuous columns variable correlation
data1 = data[['age', 'trestbps', 'chol', 'thalach', 'oldpeak']]
sns.pairplot(data1)
plt.show()
```



Outlier Detection

Since the dataset is not large, we cannot discard the outliers. We will treat the outliers as potential observations.

```
[13]: fig, ax = plt.subplots(figsize=(12,10))
      sns.boxplot(data=data, ax=ax)
      plt.show()
```



Handling Imbalance

Imbalance in a dataset leads to inaccuracy and high precision, recall scores. There are certain resampling techniques such as undersampling and oversampling to handle these issues.

Considering our dataset, the response variable target has two outcomes “Patients with Heart Disease” and “Patients without Heart Disease”. Let us now observe their distribution in the dataset.


```
[14]: data["target"].value_counts()
```

```
[14]: 1    165
      0    138
      Name: target, dtype: int64
```

```
[15]: # To study ST depression relationship with heart disease in reference to male
      ↪ and female
      plt.figure(figsize=(12, 10))
      sns.boxplot(x="target", y="oldpeak", hue="sex", data=data);
      plt.title('ST depression vs. Heart Disease')
      plt.xlabel('Heart Disease')
```



```
plt.ylabel('ST depression')
plt.show()
```




Heart disease Positive patients exhibit a lowered median for ST depression level, while negative patients have higher levels. No much differences between male & female target outcomes, expect for the fact that males have slightly larger ranges of ST Depression

Relationship between age and heart disease

```
[16]: sns.boxplot(x='target', y='age', data=data)

# Add labels and title
plt.title('Relationship between Age and Heart Disease')
plt.xlabel('Heart Disease (0 = No, 1 = Yes)')
plt.ylabel('Age')

# Show the plot
plt.show()
```



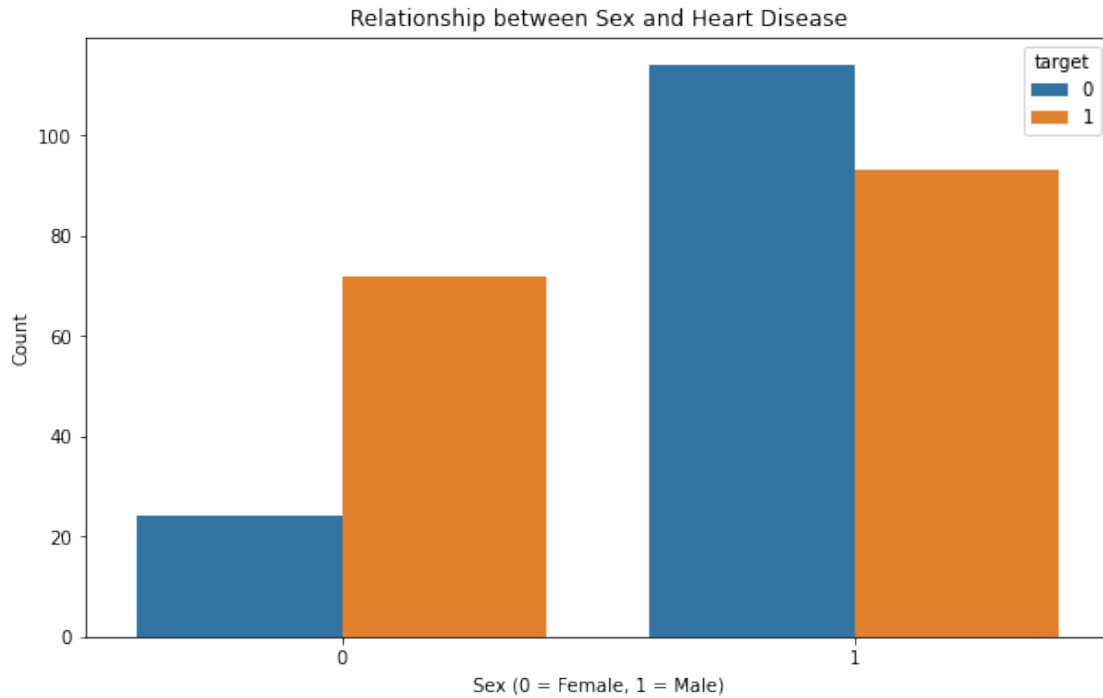
Relationship between gender and heart disease

```
[17]: plt.figure(figsize=(10, 6))

# Count plot to show the distribution of sex for each target class
sns.countplot(x='sex', hue='target', data=data)

# Add labels and title
plt.title('Relationship between Sex and Heart Disease')
plt.xlabel('Sex (0 = Female, 1 = Male)')
plt.ylabel('Count')

# Show the plot
plt.show()
```



The above chart depicts Male are more prone to heart diseases than females

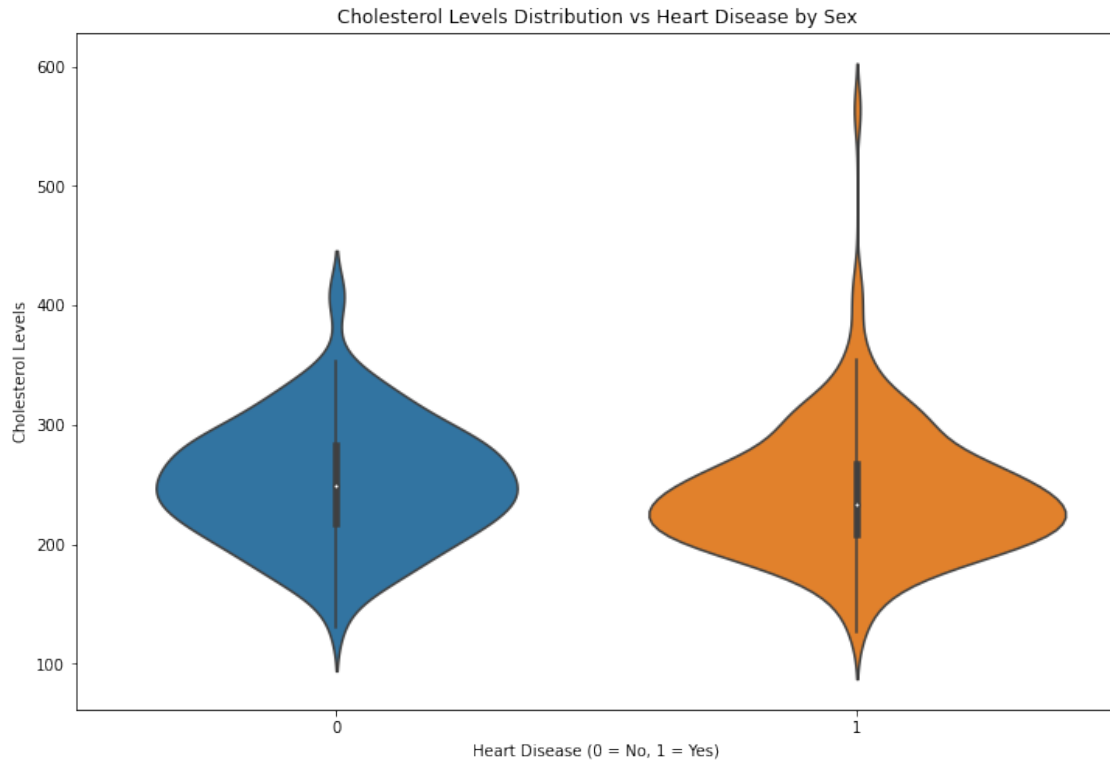
Relationship between cholestrol levels and heart disease

```
[18]: plt.figure(figsize=(12, 8))

# Violin plot to show cholesterol distribution for each target class and by sex
sns.violinplot(x='target', y='chol', data=data)

# Add labels and title
plt.title('Cholesterol Levels Distribution vs Heart Disease by Sex')
plt.xlabel('Heart Disease (0 = No, 1 = Yes)')
plt.ylabel('Cholesterol Levels')

plt.show()
```

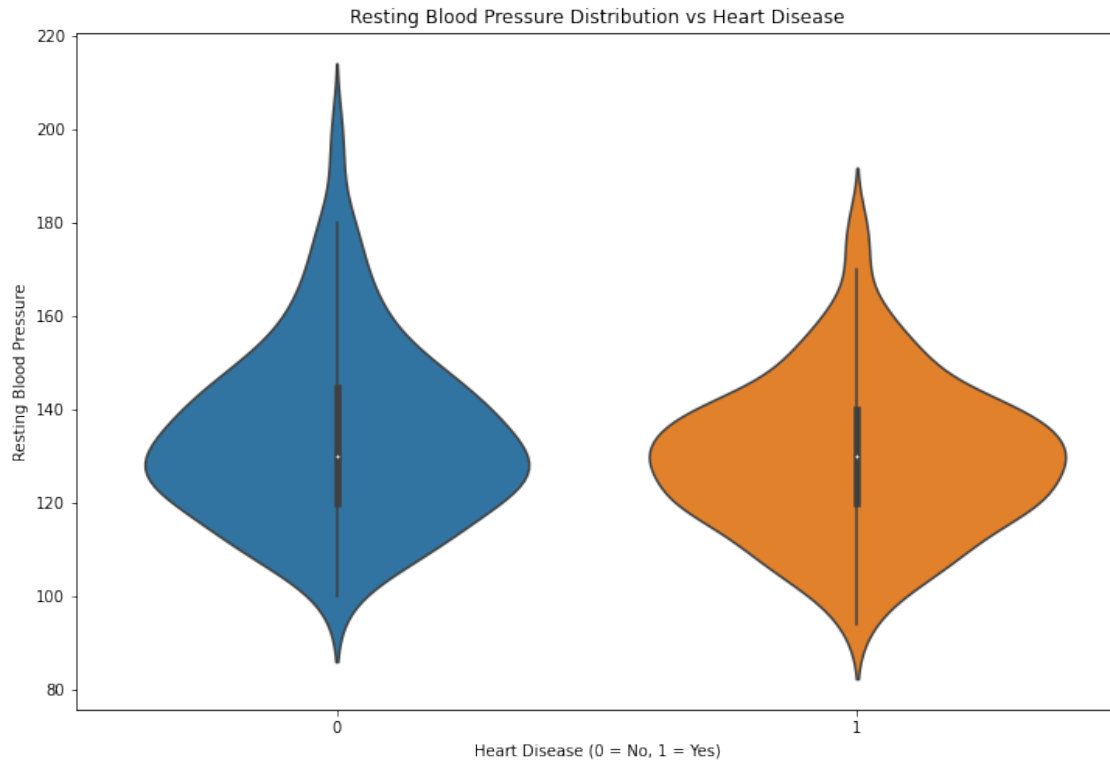


```
[19]: plt.figure(figsize=(12, 8))


# Violin plot to show resting blood pressure distribution for heart disease ↵
      ↪status
sns.violinplot(x='target', y='trestbps', data=data)

# Add labels and title
plt.title('Resting Blood Pressure Distribution vs Heart Disease')
plt.xlabel('Heart Disease (0 = No, 1 = Yes)')
plt.ylabel('Resting Blood Pressure')

plt.show()
```




```
[20]: # Boxplot for Maximum Heart Rate vs Heart Disease
plt.figure(figsize=(12, 6))
sns.boxplot(x='target', y='thalach', data=data)
plt.title('Maximum Heart Rate vs Heart Disease')
plt.xlabel('Heart Disease (0 = No, 1 = Yes)')
plt.ylabel('Maximum Heart Rate Achieved')
plt.show()
```



People with heart disease had highest heart rate

```
[21]: # Count Plot for Exercise-Induced Angina vs Heart Disease
plt.figure(figsize=(10, 6))
sns.countplot(x='exang', hue='target', data=data)
plt.title('Exercise-Induced Angina vs Heart Disease')
plt.xlabel('Exercise-Induced Angina (0 = No, 1 = Yes)')
plt.ylabel('Count')
plt.show()
```



```
[22]: from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LogisticRegression
```

```
[23]: # Defining X and Y for modeling
      X = data.drop("target",axis=1)
      y = data["target"]
      X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.
      ↪20,stratify=y,random_state=7)
```

```
[24]: #Normalize: Standardizing the data will transform the data so that its
      ↪distribution will have a mean of 0 and a standard deviation of 1
      from sklearn.preprocessing import StandardScaler
      sc = StandardScaler()
      X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)
```

```
[25]: lr = LogisticRegression()
      lr.fit(X_train, y_train)
```

```
[25]: LogisticRegression()
```

```
[ ]: pred = lr.predict(X_test)
      r2_test = lr.score(X_test,y_test)
      print(f"R-squared on testing data: {r2_test}")
```

```
r2_train= lr.score(X_train,y_train)
print(f"R-squared on training data: {r2_train}")
```

```
[ ]: from sklearn.metrics import accuracy_score, confusion_matrix, \
      classification_report
```

```
[ ]: # Accuracy on Test data
      accuracy_score(y_test, pred)
```

```
[ ]: # Accuracy on Train data
      accuracy_score(y_train, lr.predict(X_train))
```

Building a predictive system

```
[ ]: import warnings
      in_data = (57,0,0,140,241,0,1,123,1,0.2,1,0,3)

      # Changing the input data into a numpy array
      in_data_as_numpy_array = np.array(in_data)

      # Reshaping the numpy array as we predict it
      in_data_reshape = in_data_as_numpy_array.reshape(1,-1)
      pred = lr.predict(in_data_reshape)
      print(pred)

      if(pred[0] == 0):
          print('The person does not have heart disease.')
      else:
          print('The person has heart disease.')
```

```
[ ]:
```